

Physical Design


- Design so far
 - E/R modelling helps find the requirements of a database
 - Normalisation helps to refine a design by removing data redundancy
- Physical design
 - Concerned with storing and accessing the data
 - How to deal with media failures
 - How to access information efficiently

RAID Arrays

- RAID - redundant array of independent (inexpensive) disks
 - Storing information across more than one physical disk
 - Speed - can access more than one disk
 - Robustness - if one disk fails it is OK
- RAID techniques
 - Mirroring - multiple copies of a file are stored on separate disks
 - Striping - parts of a file are stored on each disk
 - Different levels (RAID 0, RAID 1...)


RAID Level 0

- Files are split across several disks
 - For a system with n disks, each file is split into n parts, one part stored on each disk
 - Improves speed, but no redundancy



RAID Level 1

- As RAID 0 but with redundancy
 - Files are split over multiple disks
 - Each disk is mirrored
 - For n disks, split files into $n/2$ parts, each stored on 2 disks
 - Improves speed, has redundancy, but needs lots of disks



Parity

- Parity - for a set of data in binary form we count the number of 1s for each bit across the data
- If this is even the parity is 0, if odd then it is 1

1 0 1 1 0 0 1 1

0 0 1 1 0 0 1 1

1 0 1 0 1 0 0 1

0 1 1 0 1 1 1 0

0 1 0 0 0 1 1 1

Recovery With Parity

- If one of our pieces of data is lost we can recover it
 - Just compute it as the parity of the remaining data and our original parity information

1 0 1 1 0 0 1 1


0 0 1 1 0 0 1 1

0 1 1 0 1 1 1 0

0 1 0 0 0 1 1 1

RAID Level 3

- Data is striped over disks, and a parity disk for redundancy
 - For n disks, we split the data in $n-1$ parts
 - Each part is stored on a disk
 - The final disk stores parity information



Other RAID Issues

- Other RAID levels consider
 - How to split data between disks
 - Whether to store parity information on one disk, or spread across several
 - How to deal with multiple disk failures

Indexes (discussed)

Query Processing

- Once a database is designed and made we can query it
 - A query language (such as SQL) is used to do this
 - The query goes through several stages to be executed
- Three main stages
 - Parsing and translation - the query is put into an internal form
 - Optimisation - changes are made for efficiency
 - Evaluation - the optimised query is applied to the DB

Parsing and Translation

- SQL is a good language for people
 - It is quite high level
 - It is non-procedural
- Given an SQL statement we want to find an equivalent relational algebra expression
- This expression may be represented as a tree - the query tree

Some Relational Operators

- Product \times
 - Product finds all the combinations of one tuple from each of two relations
 - $R1 \times R2$ is equivalent to

```
SELECT *  
FROM R1, R2
```
- Selection σ
 - Selection finds all those rows where some condition is true
 - $\sigma_{\text{cond}} R$ is equivalent to

```
SELECT *  
FROM R  
WHERE <cond>
```

Some Relational Operators

- Projection π
 - Projection chooses a set of attributes from a relation, removing any others
- $\pi_{A1,A2,\dots} R$ is equivalent to

```
SELECT
    A1, A2, ...
FROM R
```
- Projection, selection and product are enough to express queries of the form

```
SELECT <cols>
FROM <table>
WHERE <cond>
```

SQL → Relational Algebra

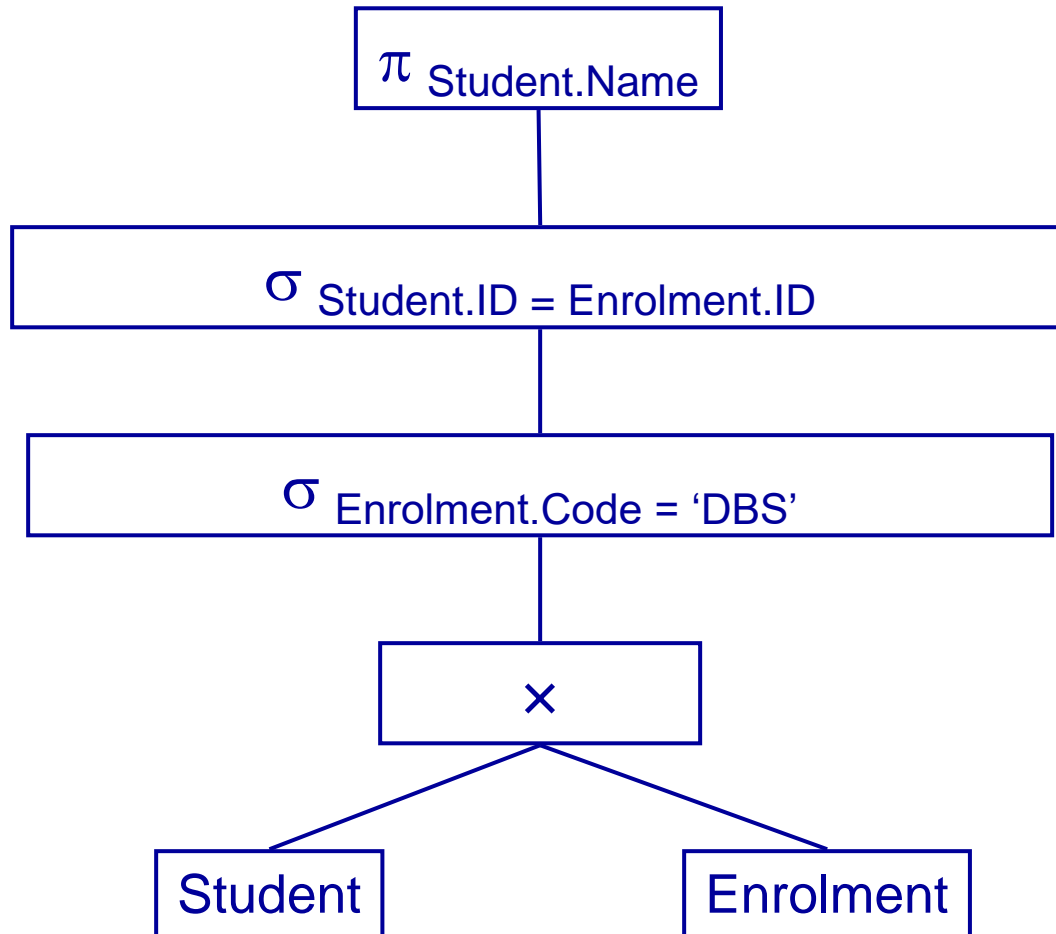
- SQL statement

```
SELECT Student.Name
FROM Student,
      Enrolment
WHERE
      Student.ID =
      Enrolment.ID
AND
      Enrolment.Code =
      'DBS'
```

- Relational Algebra

- Take the product of Student and Enrolment
- select tuples where the IDs are the same and the Code is DBS
- project over Student.Name

Query Tree



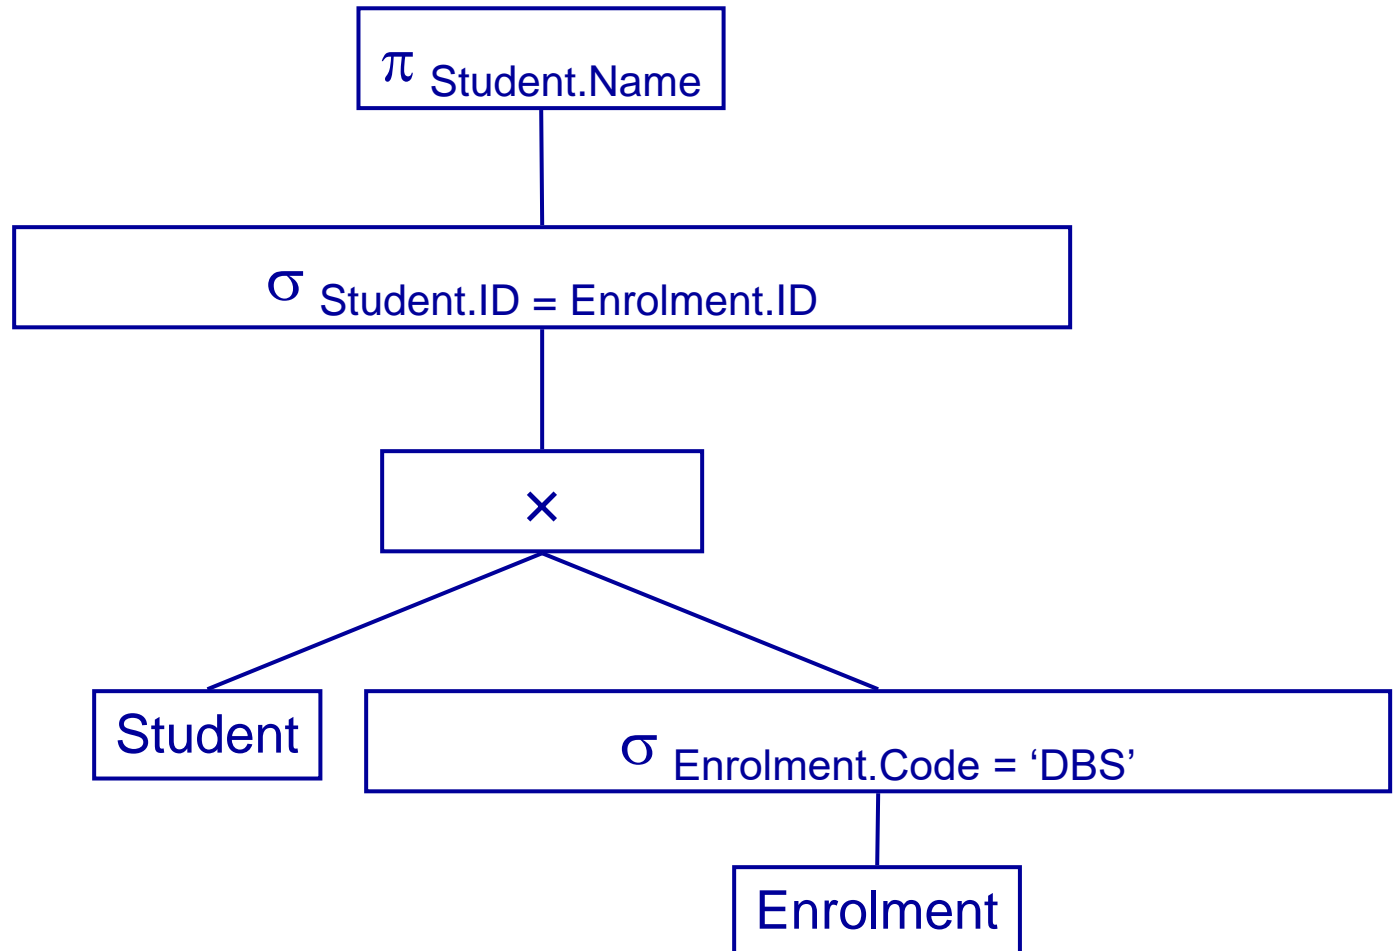
Optimisation

- There are often many ways to express the same query
- Some of these will be more efficient than others
- Need to find a good version
- Many ways to optimise queries
 - Changing the query tree to an equivalent but more efficient one
 - Choosing efficient implementations of each operator

Optimisation Example

- In our query tree before we have the steps
 - Take the product of Student and Enrolment
 - Then select those entries where the Enrolment.Code equals 'DBS'
- This is equivalent to
 - selecting those Enrolment entries with Code = 'DBS'
 - Then taking the product of the result of the selection operator with Student

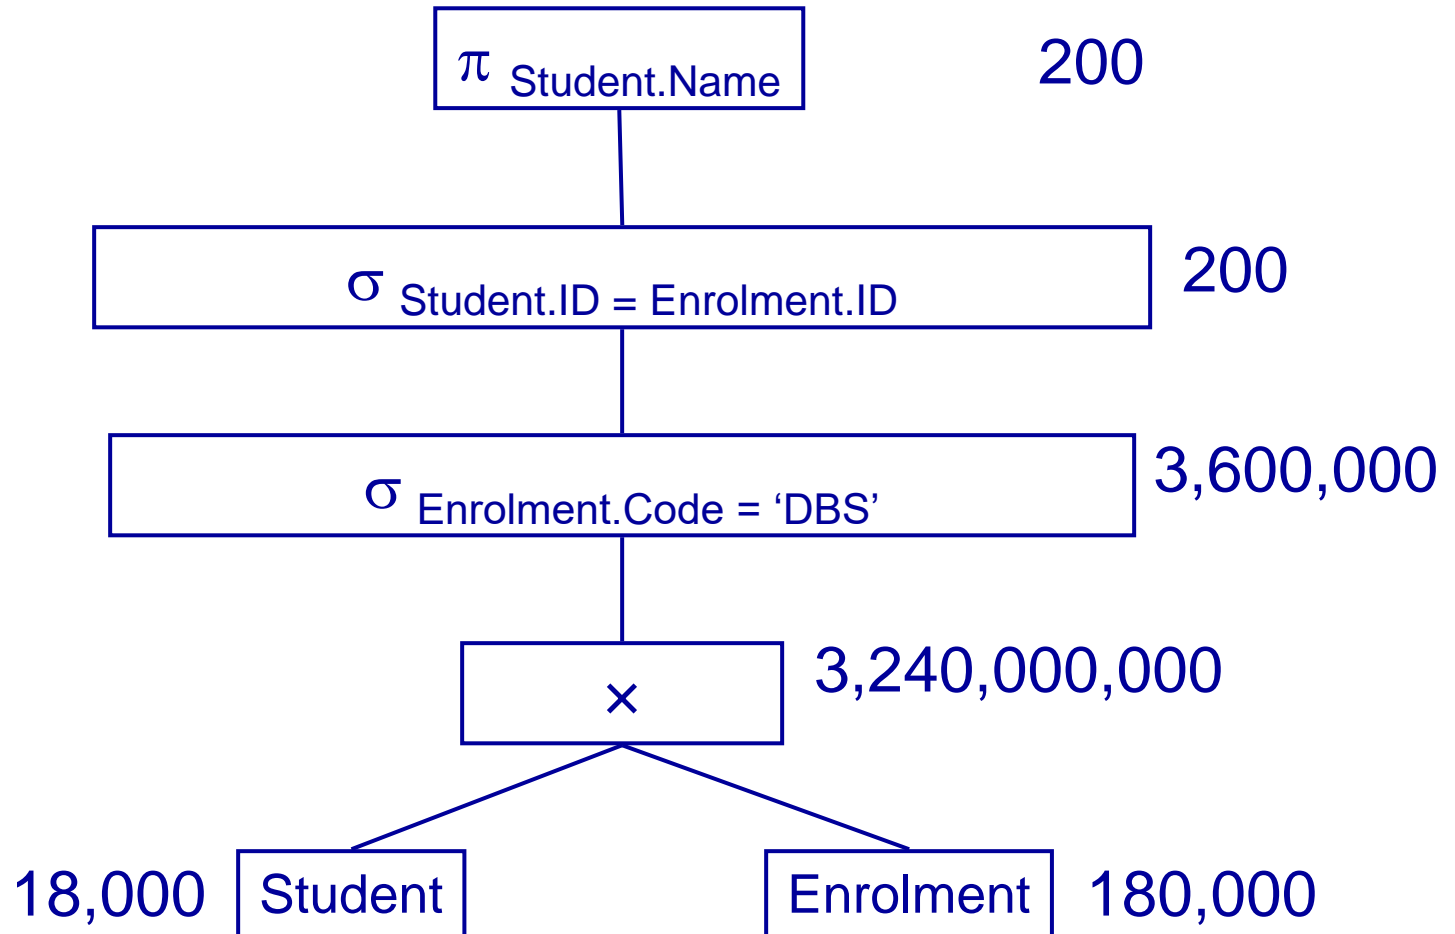
Optimised Query Tree



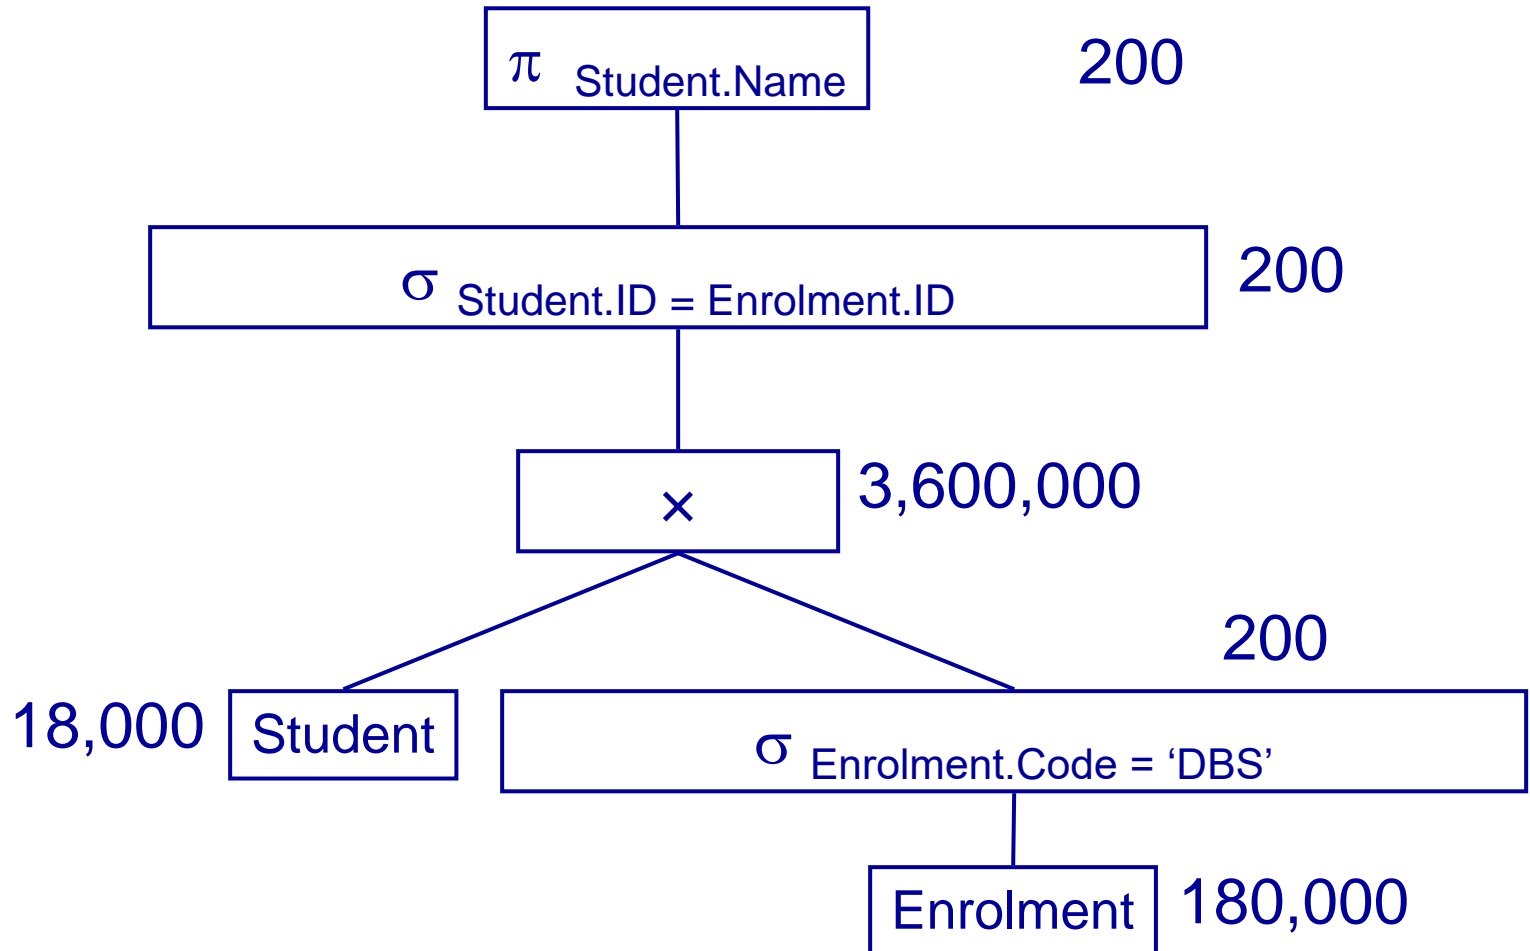
Optimisation Example

- To see the benefit of this, consider the following statistics
 - Nottingham has around 18,000 full time students
 - Each student is enrolled in at about 10 modules
 - Only 200 take DBS
- From these statistics we can compute the sizes of the relations produced by each operator in our query trees

Original Query Tree



Optimised Query Tree



Optimisation Example

- The original query tree produces an intermediate result with 3,240,000,000 entries
- The optimised version at worst has 3,600,000
- A big improvement!
- There is much more to optimisation
 - In the example, the product and the second selection can be combined and implemented efficiently to avoid generating all Student-Enrolment combinations

Optimisation Example

- If we have an index on Student.ID we can find a student from their ID with a binary search
- For 18,000 students, this will take at most 15 operations
- For each Enrolment entry with Code 'DBS' we find the corresponding Student from the ID
- $200 \times 15 = 3,000$ operations to do *both* the product and the selection.